

Lamar Holding — Round 2 schedule

This page fills every [TIME] / [DATE] placeholder in the brief that follows.

Deadline: 17:00 on Wednesday 23 September 2026, Riyadh time.

Working session: Friday 25 September 2026, Riyadh time.

Ninety minutes, by video, screen shared. Your exact start time is in the covering email.

Where to send: hiringforhani@gmail.com

Please confirm these dates by Friday 18 September 2026.

Everything else in the brief stands as written.

PPP AI Partner

Round 2 — Working Assessment

Candidate brief

Lamar Holding W.L.L. · Development Division · PPP team

Read this page before you start.

This is a five-day exercise followed by a ninety-minute working session. It is the last step before an offer.

It is not a product brief. We are not asking you to design a tool for us, pitch a platform, or show us what is possible. We are asking you to be right about a specific set of documents, and to show us how you got there.

The assignment

In one paragraph

You are given six documents from a public-private partnership tender pack for a sewage treatment plant: Instructions to Bidders, Technical Requirements, Form Sheets, an extract of the draft Project Agreement, and two addenda. Build a working system that reads them and produces five artefacts — a compliance register, an addendum reconciliation, a one-page list of what would disqualify a bid, a bid programme derived from that register, and a work log. You have five days. Then we meet for ninety minutes, we hand you a third addendum you have not seen, and you make your system absorb it while we watch.

What you are given	Six PDF documents, thirty-two pages, attached to the email that sent you this brief. Nothing else. The tender is fictional and was written for this exercise.
What you build	A working system — scripts, a repository, whatever you normally build with — that turns those documents into the five artefacts in Section 3. Not a prototype of a product. A thing that produces correct output.
What you send back	The five artefacts (A1 to A5 in Section 3) and the repository, to the address below, by the deadline below. A single archive is fine.
How long	Five days. We are not counting hours. If you finish in two, send it in two.
Deadline	[TIME] on [DATE], Riyadh time.
Working session	[TIME] on [DATE], Riyadh time. Ninety minutes, by video, screen shared, you driving. Have your development environment open and running.
Where to send it	hiringforhani@gmail.com. Send questions to the same address at any time.

Sections 1 to 7 explain the reasoning, the detail of each artefact, what not to build, how the working session runs and how we score it. Read them before you start. If you remember one thing from this document, make it this: **we are not judging what you build. We are judging whether it is right, and whether you can change it in front of us.**

1. Why this exercise looks the way it does

This round is built differently from the first, and it is worth knowing why before you start. The seat exists because we want AI used hard, every day, next to our bid team. That is precisely why a demonstration tells us so little. What we need to know is whether a system is right about a specific set of documents, and whether the person who built it can change it under pressure.

You will work on a tender document pack that we wrote. It has never been published, it is not in any training set, and we know every correct answer in it before you begin. Your output is marked against that key. A polished interface earns nothing. Being right earns everything.

What this means in practice

Use whatever models, tools and assistance you like. We expect you to. We ask only that you log it, and that you can defend every number your system produces.

If your system cannot determine something, say so. An honest 'I could not resolve this' scores better with us than a confident answer that turns out to be invented. One wrong requirement in a live bid pack can disqualify a bid worth hundreds of millions.

2. The documents

Six documents, thirty-two pages, from a fictional Saudi PPP tender for a sewage treatment plant procured on a BOOT basis. Instructions to Bidders, Technical Requirements, Form Sheets, a draft Project Agreement extract, and two addenda.

The pack is small on purpose. Nothing in it is hidden behind volume. Everything we will ask you about can be resolved from those pages by a careful person in an afternoon. Your software is there to make that repeatable and auditable, not to make it possible.

3. What you deliver

Five artefacts. Nothing else.

	Artefact	What it must contain
A1	Obligations and compliance register	One row per requirement. A requirement identifier; the verbatim source text; the document, clause and page it came from; whether it is pass/fail or scored; the discipline that owns it; the evidence needed to prove compliance; its status after each addendum; and your confidence. Machine-readable — CSV or JSON — plus whatever interface you prefer to work in.
A2	Addendum reconciliation	What each addendum changed, added and deleted. Which register rows move as a result. Which earlier answers are now wrong. We must be able to trace any row back through the chain to the original document.
A3	One page: what would disqualify us	Every requirement whose breach would put a bid out — the ones the documents themselves say cause rejection, disqualification or non-responsiveness, not every obligation in the pack — with your confidence in each. And, explicitly, the list of things your system could not resolve, and why. We read this page first.

	Artefact	What it must contain
A5	Bid programme and submission marshalling plan	Generated by your system from A1 — not drawn by hand. Two parts. First, the programme: the activities needed to get a compliant bid out of the door, with durations, dependencies, owners by discipline, and resources. Every activity must carry the requirement ID or IDs from A1 that it discharges, and every date must be computed backwards from the deadline, not typed in. Second, the marshalling plan: the last stretch before submission, listing every physical and documentary item the pack requires, who issues it, how long that takes, and the drop-dead date to start it. Output the plan as data, CSV or JSON. If it reads better as a chart, show us the chart too — but the chart is not the deliverable.
A4	Work log	The repository with its real commit history. The prompts and model calls you used. And a short written note of every place your system was wrong while you were building it, and how you caught it.

A4 is not a formality and it is not a trap. We want to see how you work with a model, including where it misled you. A work log that records no errors tells us you were not checking.

4. What not to build

These are instructions, not preferences. Submitting one of these in place of the five artefacts ends the assessment — see Section 6.

- **No chat interface.** Nobody on our bid team wants to interrogate a tender by conversation under deadline.
- **No multi-agent framing.** Do not name components after departments.
- **No dashboard as a product,** no landing page, no logo, no product name, no roadmap, no pricing. To be exact, because A5 asks for a programme: a view of your own plan is welcome and a Gantt chart is a perfectly good way to show one. What we do not want is a dashboard built to impress — tiles, counters, a product shell — wrapped around output we cannot trace.
- **No strategy document.** We already have one.
- **Do not tell us what else you could build.** Build the five artefacts and make them correct.

If you find yourself writing a paragraph that begins ‘this could be extended to’, delete it and go back to the register.

5. The working session

Ninety minutes, screen share, you drive. Three of us will be on the call. It is a working session, not a presentation — there are no slides.

Minutes	What happens
0 – 10	We pick three rows from your register at random and ask you to show us the clause and page each one came from.
10 – 40	We hand you a third addendum you have not seen. You ingest it, tell us what it broke, and replan. At the keyboard, live.
40 – 58	We name something your system got wrong. You change the system, in front of us, so that it catches that kind of error in future — not just that one instance.
58 – 70	You explain one part of what you built to a bid coordinator with no technical background, well enough that they could use it on Monday.
70 – 85	Cost per bid if we run eight bids at once. Where it falls over. What you would refuse to automate, and why.
85 – 90	Your questions for us.

Come ready to write code on the call.

Your own machine, your own editor, your own environment, already running. We are not testing recall and we are not going to ask you to write an algorithm on a whiteboard. We want to watch you change a system you built, with us watching, on a document you have not seen.

If you cannot do that, this seat is the wrong seat, and it is better for both of us to find that out now.

6. How we judge it

One hundred marks, scored independently by each of us before we speak to one another.

Marks	Dimension
15	Accuracy against our key — did you find what is actually there, and nothing that is not
10	Addendum reconciliation — is the final state of the pack correct and is the chain auditable
30	The live change — the unseen addendum, the replan and the fix, at the keyboard
15	The bid programme — is it derived from your register, and does it survive the lead times the documents actually impose
10	Calibration — do you know what you do not know, and did you say so
10	Teach-back — can a non-technical colleague operate what you built
10	Cost, scale, and where you draw the line on automation

The weighting is deliberate and worth understanding before you plan your five days. We ran the pack through a capable model on its own, with no human involved. It got nearly all of the first two rows right. So those two rows tell us very little about you, and we have weighted them accordingly. The marks are concentrated where a model cannot help you: changing a system you built, under time pressure, on a document you have never seen, while three people watch.

Five things end the assessment regardless of score

Stated here in full so that none of them is a surprise.

- **Asserting a requirement that is not in the pack.** If your register states a requirement of this tender that no document supports, with no source and no flag, that ends it. A tool that invents a requirement under deadline is worse than no tool, because a person will believe it.
- **Being unable to show where a row came from.** We will pick rows at random and ask for the clause and the page, and we expect an answer inside a minute.
- **Being unable to make a code change during the session.**
- **Submitting something other than the five artefacts** — a dashboard, a chat tool, a deck, a pitch. Section 4 is not a style note.
- **Letting the tool decide a commercial or legal question** that a person has to own. Some things in this pack cannot be settled by software, and one of them has no correct answer at all. Surfacing those to a human is the right result. A confident machine answer to them is the wrong one.

7. Terms

Item	Detail
Time	Five days. We are not measuring hours and we are not counting them. If you finish in two, submit in two and say so.
Questions	Ask them. Send anything to hiringforhani@gmail.com at any point. We answer every question to all candidates, so asking costs you nothing and tells us something.
Ownership	What you build stays yours. If you join, we will want to build on it, and we will agree that separately.
Confidentiality	The pack is fictional and contains no Lamar confidential material. What you learn about how we bid is confidential.
Other offers	If you are in another process, tell us. We will work to your timeline rather than lose you to a scheduling problem.

One last thing.

We are not looking for the most impressive thing anyone can build in five days. We are looking for the person we would want sitting next to the bid team the week a tender lands, when the addendum arrives at four in the afternoon and something that was true this morning is no longer true.

Build for that day.